


## Insights on Model Performance



- The model with Dropout Regularization performed well, indicating reduced overfitting.
- Increasing hidden units to 128 slightly improved accuracy compared to 32 units, showing the importance of capacity.
- Models using tanh activation and MSE loss showed lower accuracy, suggesting they are less effective for binary classification

---

### IMDB Sentiment Analysis Model Performance Analysis

#### Model Variations and Results

Our experiments tested several architectural and hyperparameter variations on the IMDB sentiment analysis task, with the following key findings:

## Model Architecture Impact

### 1. Layer Depth:

- Single hidden layer (64 units) achieved 86.048% accuracy
- Three hidden layers performed slightly worse at 84.764%
- This suggests that for this particular task, deeper architectures don't necessarily improve performance
- The simpler architecture may be sufficient for capturing the necessary sentiment patterns

### 2. Hidden Unit Variations:

- 32 units: 86.032% accuracy
- 64 units: 86.048% accuracy
- 128 units: 86.524% accuracy
- The trend shows a slight improvement with increased units
- The marginal gains diminish as we add more units
- The 128-unit model performed best overall, suggesting this capacity level is optimal for the task

## Training Optimizations

### 1. Activation and Loss Function:

- The tanh activation with MSE loss performed notably worse (82.756%)

- This validates modern best practices of using ReLU and binary cross-entropy for binary classification tasks
- The significant performance drop ( $\approx 4\%$  decrease) demonstrates the importance of appropriate activation/loss function selection

## 2. Regularization Impact:

- Dropout regularization (85.856%) performed nearly as well as the best model
- The small gap between validation and test accuracy suggests effective prevention of overfitting
- This indicates that dropout is a valuable addition to the model architecture

## Key Takeaways

### 1. Model Complexity vs. Performance:

- Simpler architectures performed surprisingly well
- Adding more layers didn't improve performance
- This suggests the sentiment classification task may not require deep architectural complexity

### 2. Optimal Configuration:

- Best performance: 128 units with ReLU activation and binary cross-entropy loss

- Dropout provides good regularization without significant performance penalty
- The modern standard of ReLU activation significantly outperforms traditional tanh

### 3. Practical Implications:

- For similar text classification tasks, starting with a single hidden layer and ReLU activation is recommended
- Increasing model width (units) is more beneficial than increasing depth (layers)
- Dropout should be considered as a standard addition to prevent overfitting

### Future Recommendations

#### 1. Consider experimenting with:

- Different dropout rates to find optimal regularization
- Embedding layer dimensionality
- Additional regularization techniques (L1/L2)
- Different optimizers beyond Adam

#### 2. Performance improvements might be achieved through:

- Text preprocessing optimizations
- Longer training periods

- Learning rate scheduling
- Ensemble methods combining multiple model variants

#### Final Recommendation:

- Model: Single hidden layer with 128 units, ReLU activation, and binary cross-entropy loss
- Reasoning:
  - Achieved the highest test accuracy.
  - Simpler architecture (fewer layers) performed better than deeper models.
  - ReLU activation outperformed tanh.
  - Dropout helped prevent overfitting but didn't outperform the 128-unit model significantly.