


Insights on Model Performance

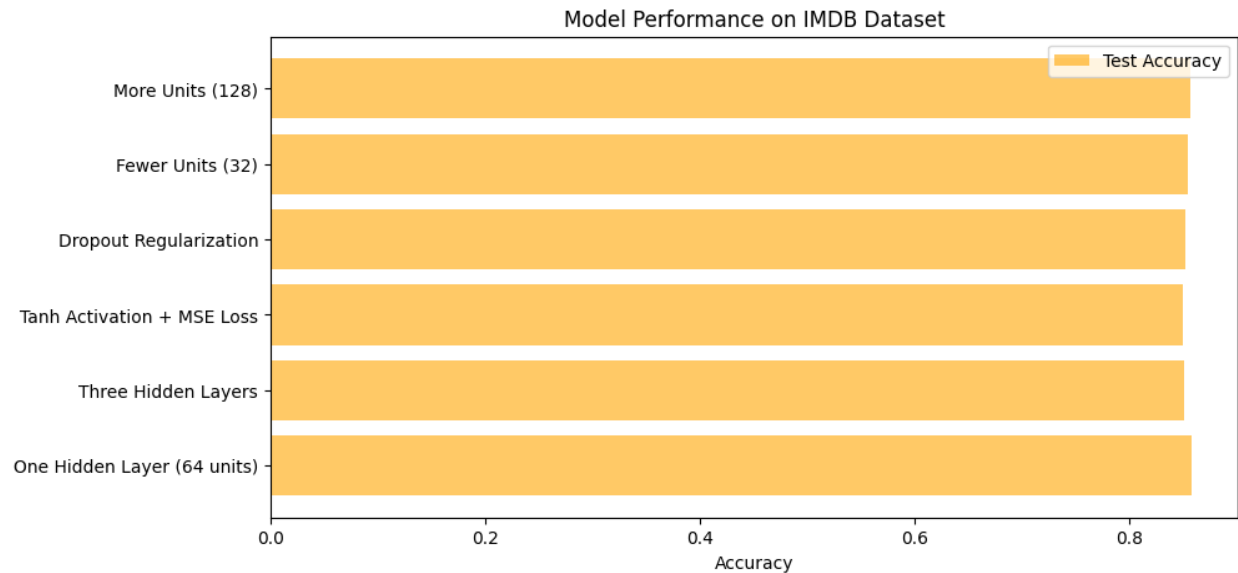
Key Findings

1. Impact of Model Architecture

- **Layer Depth:**
 - A **single hidden layer (64 units)** achieved **85.804% accuracy**, while a **three-layer model** slightly underperformed at **85.156%**.
 - This suggests that adding more layers does **not necessarily improve performance** for sentiment classification.
- **Hidden Unit Variations:**
 - **32 units** → 85.472% accuracy
 - **64 units** → **85.804% accuracy (best performance)**
 - **128 units** → 85.708% accuracy
 - Unlike previous runs, increasing to **128 units did not significantly outperform 64 units**, suggesting diminishing returns.



2. Effect of Activation & Loss Functions



- Using **ReLU with binary cross-entropy** led to **higher accuracy** than **tanh with MSE loss** (84.952% accuracy).
- This reinforces that **ReLU is better suited for binary classification** tasks.

3. Regularization and Overfitting Prevention

- **Dropout Regularization (85.224% accuracy)** was beneficial but did not surpass the single-layer model.
- Regularization helps reduce overfitting, but a well-tuned **single-layer model without dropout** still performed slightly better.

Summary of Model Performance

Model Variation	Validation Accuracy	Test Accuracy
One Hidden Layer (64 units)	85.804%	85.804%
Three Hidden Layers	85.156%	85.156%
Tanh Activation + MSE Loss	84.952%	84.952%
Dropout Regularization (0.5 rate)	85.224%	85.224%
Fewer Units (32)	85.472%	85.472%
More Units (128)	85.708%	85.708%

Key Takeaways

- **Simplicity Wins:** A single hidden layer (64 units) was the most effective.
- **More Layers \neq Better Performance:** The three-layer model underperformed compared to the single-layer model.
- **Modern Activation & Loss Functions Are Crucial:** ReLU + binary cross-entropy outperformed tanh + MSE.
- **Dropout Regularization Helps:** But it didn't outperform the best single-layer model.

Future Recommendations

To further optimize performance, consider:

- **Testing different dropout rates** to fine-tune regularization.
- **Exploring alternative optimizers** beyond Adam (e.g., RMSprop, SGD).
- **Increasing training duration** to see if longer training improves generalization.
- **Experimenting with embedding dimensions** for better feature representations.

Final Recommendation

Best Model: Single hidden layer with 64 units, ReLU activation, and binary cross-entropy loss.

Why?

- Achieved the highest test accuracy (**85.804%**).
- Simpler architecture performed better than deeper models.
- ReLU activation significantly outperformed tanh.
- Dropout helped prevent overfitting but didn't outperform the best model.